



Containerized GPU

Maruthi S. Inukonda
26th Jun 2018



Speaker

I am

- Maruthi Seshidhar Inukonda
- Currently Ph.D(CSE) student at CANDLE lab@CSE, IIT Hyderabad.
- M.Tech (CSE) from IIT Roorkee
- Worked in Operating Systems' Storage stack development for 13.5 years in Industry (Veritas, NetApp, EMC²)
- Using/Working on Containers from 2014.
- Member of CCICI, SNIA, IEEE, Linux, Openstack, Ceph communities



Agenda

- Virtualization
- Virtual machines
- Containers
- Physical machine vs Virtual machine vs Container
- Dockers
- GPU Dockers & Benchmarking
- Multi-tenancy & CUDA



Virtualization



What is a Virtualization? (1/4)

- A system software that provides multi-tenancy for sharing of hardware resources. It also provides isolation.

Civil engineering example

- Dormitory complex with shared resources.
Viz., lifts, generator, security staff, water pumps, television, restrooms, hall, bedroom.
- Hostel complex having multiple rooms with shared resources.
Viz., all the above but not study/bedroom.
- Apartment complex having multiple flats with shared resources.
Viz., all the above but not hall, restroom, study/bedroom, television.



What is a Virtualization? (2/4)

Computing hardware resources

- CPU, RAM, Disk, NIC.
- Accelerators (GPU)

Terminology

- Hypervisor: A software that emulates hardware/run-time environment.
- Guest: OS running in VM.
- Host: OS running on physical hardware.



What is a Virtualization? (3/4)

Types

- Full Virtualization:
 - Emulates entire computer system.
 - Runs unmodified guest OS (kernel and userspace).Eg. VirtualBox, XenServer, VMWare, Hyper-V, etc
- Para Virtualization:
 - Emulates entire computer system.
 - Runs modified guest OS kernel, and unmodified guest OS userspace.Eg. KVM



Note: We will not talk about Para virtualization



What is a Virtualization? (4/4)

Types

- Process Virtualization :
 - Emulates one programming language run-time environment.
 - Runs an application written in single programming language.Eg. Java/Python VM
- Operating System Virtualization :
 - Does not emulate any hardware. Directly uses host hardware.
 - Runs an user-space (application and libraries) of OS.Eg. Containers/Dockers



Note: We will not talk about Process virtualization



Virtual Machines



What is a Virtual Machine?

- Virtual Machine (VM) is an emulation of entire computer system, as result of full virtualization.
- Each VM runs its own operating system instance.
- Civil Engineering example :
 - Apartment complex having multiple flats.
 - Each flat has virtualized resources (generator , lift, security guard, water pump, etc.)

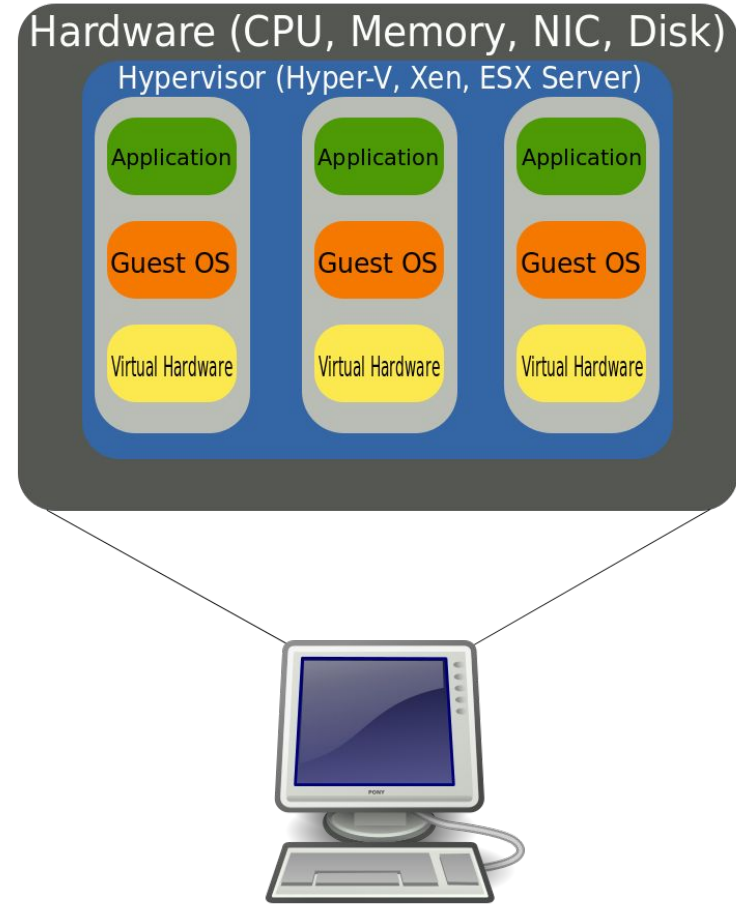


Image Courtesy: Wikipedia

Containers



What is a Container? (1/2)

- Linux Containers (LXC) is an operating-system-level virtualization method.
- For running multiple isolated Linux systems (containers) on a control host using a single Linux kernel.
- Directly runs on hardware. (No per-instruction level trapping)
- An unprivileged user on host can be privileged user on guest.
- Civil Engineering example :
 - Hostel complex having multiple rooms with shared resources.
Viz., all the above but not study/bedroom.

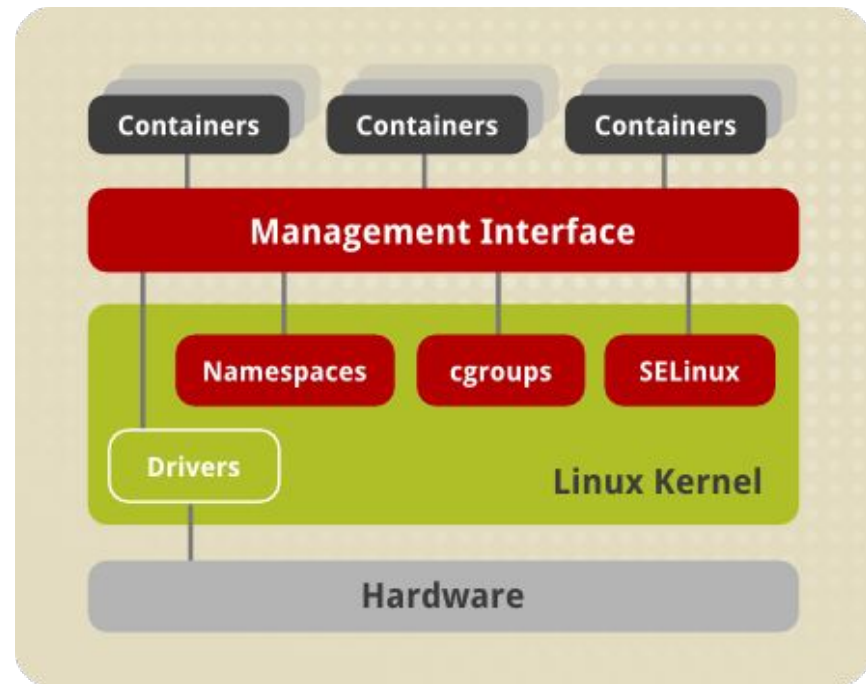


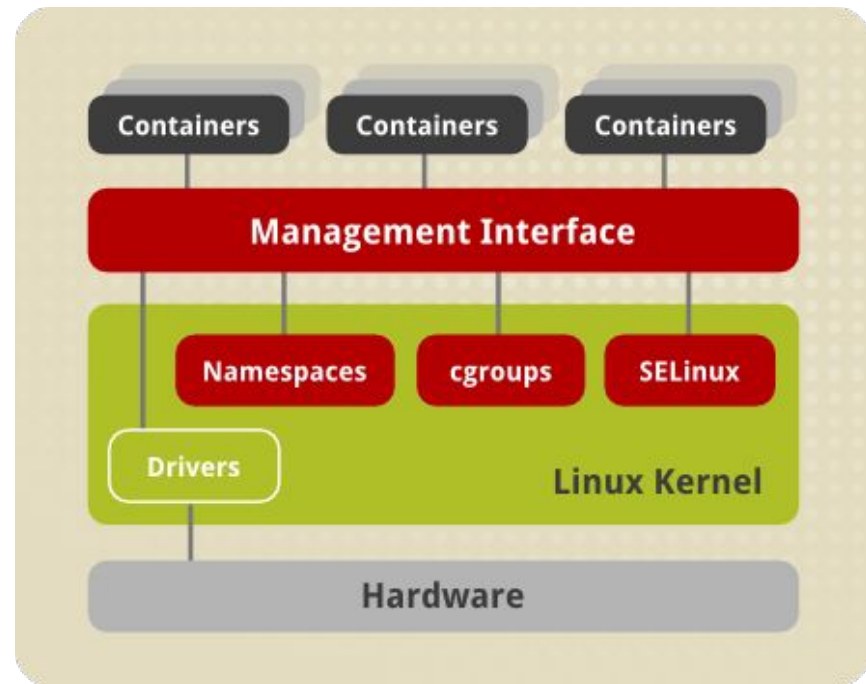
Image Courtesy: Redhat Customer Portal



What is a Container? (2/2)

It is implemented using following features in Linux

- Advanced Multi-layer Union FS (AUFS) or Overlay FS
- Kernel namespaces
- Cgroups
- Capabilities
- Netfilter, Netlink
- Bind mount
- Role-Based Access Control (RBAC)
 - Eg. SELinux, AppArmor



Physical Machine vs Virtual Machine vs Container



Physical Machine vs Virtual Machine vs Container

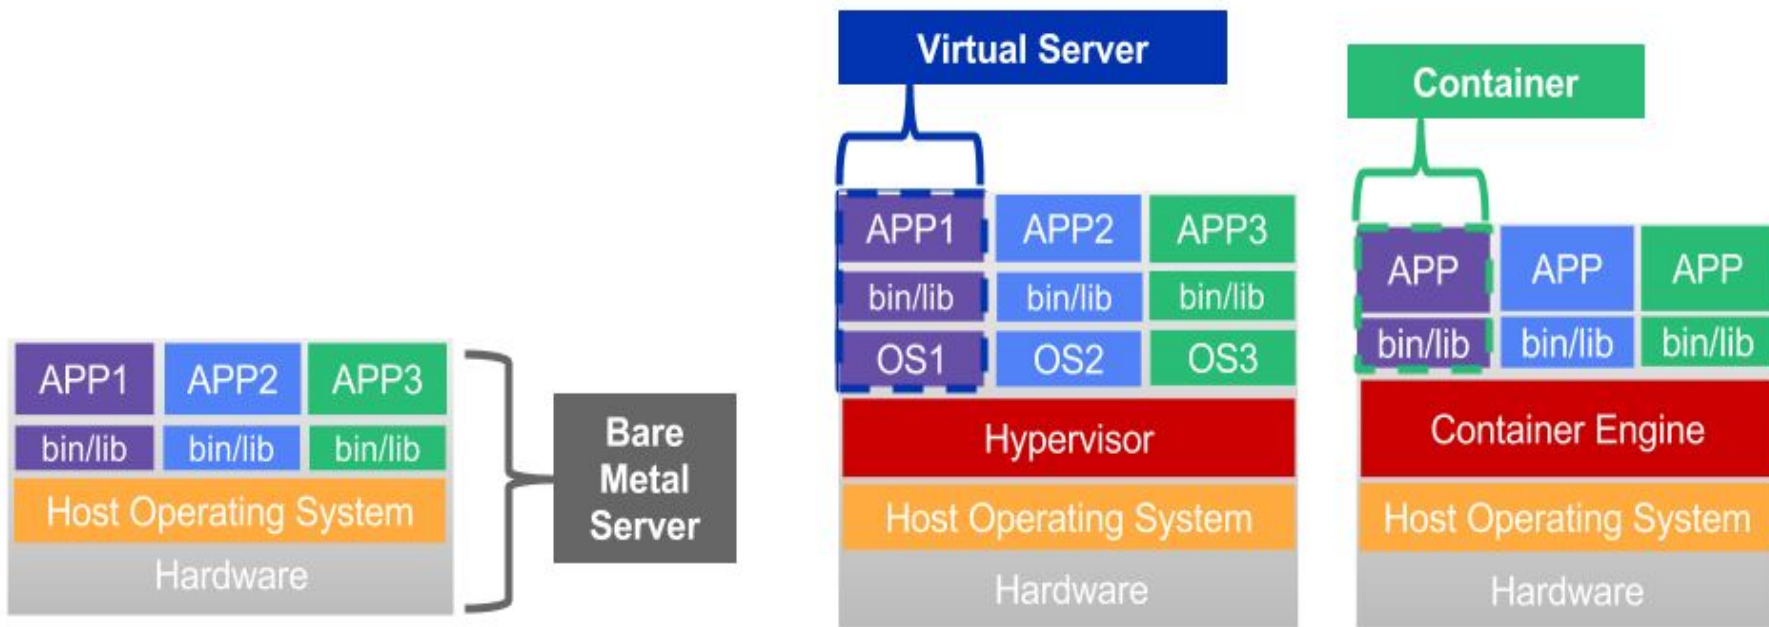


Image Courtesy: Kumulus Technologies

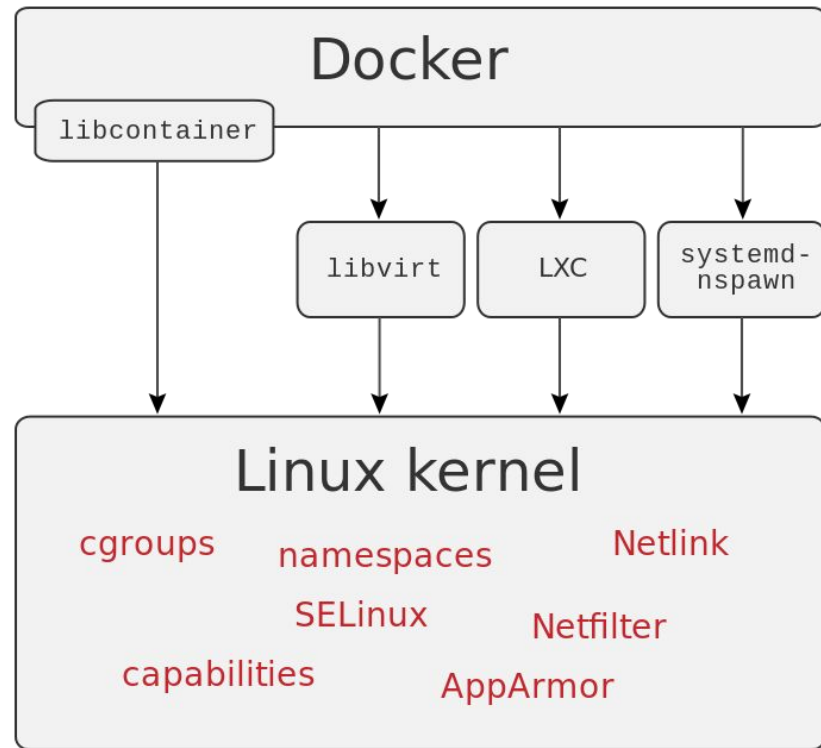


Dockers



What is a Docker?

- Docker is a company that provides software (also called Docker) that allows you to build, run and manage software containers.
- It makes container deployment and administration quite easy.
- It allows re-use of containers created by others.



Docker components

Docker has three major components

- Docker repository/hub
- Docker image
- Docker container



Image Courtesy: Wikipedia



Docker hub

Browser address bar: <https://hub.docker.com/search?isAutomated=0>

Search bar: Search


Notification: Docker Store is the new place to discover public Docker content. [Check it out →](#)

Navigation bar:

- Search: nvidia
- Dashboard
- Explore
- Organizations
- Create
- User: maruthi

Repositories (774)

Filter: All

	nvidia/cuda public	299 STARS	1M+ PULLS	DETAILS
---	-----------------------	--------------	--------------	-------------------------



Pulling images

- To pull/download docker images, use
`docker pull <image>:<tag>`

```
$ docker pull nvidia/cuda:latest
latest: Pulling from nvidia/cuda
b234f539f7a1: Pull complete
...
d056eaf3dfff4: Pull complete
7dc790b5527b: Pull complete
89bce857a556: Downloading [=====>] 463.4
MB/580.1 MB89bce857a556: Pull complete
Digest: sha256:3cdf1b5becfde8772e15dab594bc76de1cbbefd6c0f8533748854ab47e109ad1
Status: Downloaded newer image for nvidia/cuda:latest
```



Docker container life-cycle

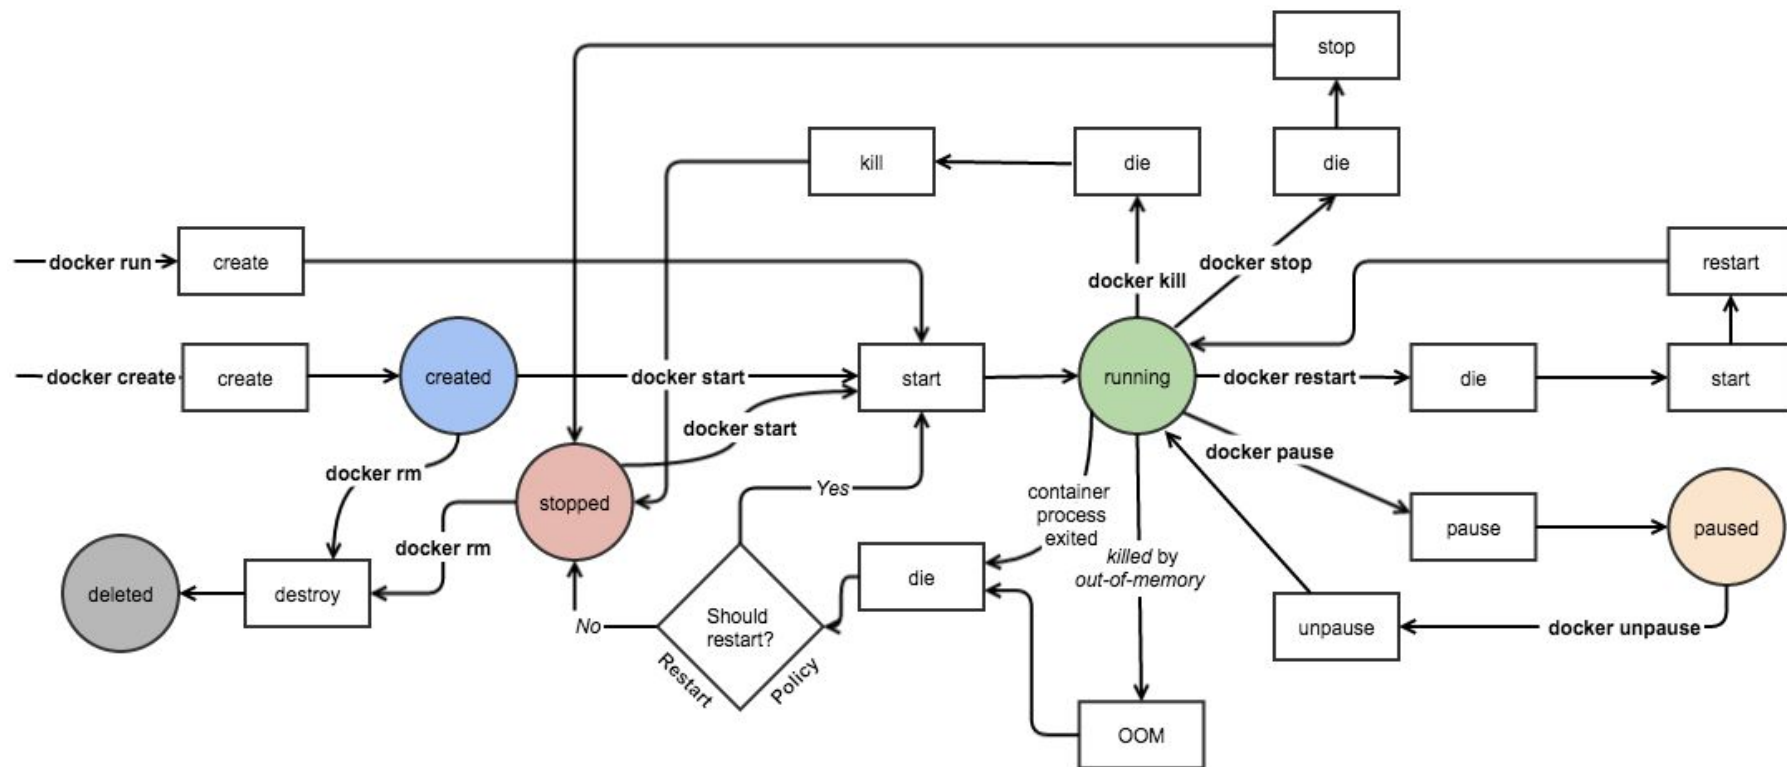


Image Courtesy: Nitin Agarwal @
medium.com

Listing images

- To list downloaded images, use `docker images`

```
$ docker images
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
nvidia/opengl	latest	b546828c2b30	12 days ago	116MB
centos	latest	49f7960eb7e4	2 weeks ago	200MB
nvidia/cuda	latest	9337ecb4311e	7 weeks ago	2.24GB
ros	kinetic-ros-base	2e1693285910	8 weeks ago	1.18GB
ubuntu	latest	452a96d81c30	8 weeks ago	79.6MB
autoware/autoware	1.6.0-kinetic	8fc60a26cc84	6 months ago	7.65GB



Running a container

- To run a container, use

```
docker run -it --name <name> <image>:<tag> <program>
```

```
$ docker run -it --name centos1 centos:latest bash
```

```
[root@e7f7395af134 /]#
```

```
[root@e7f7395af134 /]# cat /etc/redhat-release
```

```
CentOS Linux release 7.5.1804 (Core)
```

```
[root@e7f7395af134 /]#
```



Listing running containers

- To list running container, use
`docker ps`

```
$ docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
e7f7395af134	centos:latest	"bash"	8 seconds ago	Up 7 seconds		centos1



Exiting from a container after stopping

- To stop and exit from a container, use `exit`

```
[root@e7f7395af134 /]# exit
```

```
$
```



Listing all containers

- To list all (running/exited) container, use
`docker ps -a`

```
$ docker ps -a
CONTAINER ID   IMAGE          COMMAND        CREATED        STATUS        PORTS        NAMES
e7f7395af134   centos:latest  "bash"        8 seconds ago  Exited (0)    5 seconds ago
centos1
```



Exiting from a container without stopping

- To exit from a container without stopping, press
`Ctrl+p Ctrl+q`

```
[root@e7f7395af134 /]# Ctrl+p Ctrl+q      read escape sequence
```

```
$
```



Creating a container

- To create a container, use

```
docker create --name <name> -it <image>:<tag> <program>
```

```
$ docker create --name centos2 -it centos:latest bash
916dc303760db834c1aa4a9b591605ab56c91aba97bceb6fda2ca0db564f489a
```

```
$ docker ps -a
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
e7f7395af134	centos:latest	"bash"	31 seconds ago	Created		centos1
916dc303760d	centos:latest	"bash"	About a minute ago	Created		centos2



Starting a container

- To start a container, use
`docker start <name_or_id>`

```
$ docker start centos2
centos2
```

```
$ docker ps -a
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
e7f7395af134	centos:latest	"bash"	4 minutes ago	Up 3 minutes		centos1
916dc303760d	centos:latest	"bash"	3 minutes ago	Up 2 minutes		centos2



Attaching to a running container

- To start a container, use
`docker attach <name_or_id>`

```
$ docker attach centos2  
[root@916dc303760d /]#
```



GPU Dockers



nvidia-docker

- To launch a nvidia GPU docker, use

```
nvidia-docker run -it --name <name> <image>:<tag> <program>
```

```
$ docker images
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
nvidia/opengl	latest	b546828c2b30	12 days ago	116MB
nvidia/cuda	latest	9337ecb4311e	7 weeks ago	2.24GB
ros	kinetic-ros-base	2e1693285910	8 weeks ago	1.18GB
autoware/autoware	1.6.0-kinetic	8fc60a26cc84	6 months ago	7.65GB

```
$ nvidia-docker run -it --name nvcuda1 -v /mnt:/mnt nvidia/cuda:latest bash
root@ba47b6c2ad70:/#
```



Performance difference between bare metal and container

A mean pooling program is run on a matrix of 4096x4096 256 times.
An element is updated with mean of its left, right, top and bottom elements.

On bare metal:

```
$ time ./cs18resch01001_Prog
real    0m14.549s
user    0m9.768s
sys     0m4.743s
```

In container

```
# time ./cs18resch01001_Prog
real    0m14.656s
user    0m2.057s
sys     0m12.603s
```



Multi-tenancy & CUDA



Launching multiple nvidia/cuda containers

```
$ nvidia-docker run -it --name nvcuda1 -v /mnt:/mnt nvidia/cuda:latest bash  
root@ba47b6c2ad70:/#
```

```
root@ba47b6c2ad70:/# time ./cs18resch01001_Prog  
real    0m28.756s  
user    0m4.020s  
sys     0m24.702s
```

```
$ nvidia-docker run -it --name nvcuda2 -v /mnt:/mnt nvidia/cuda:latest bash  
root@22dd291ccb73:/#
```

```
root@22dd291ccb73:~# time ./cs18resch01001_Prog  
  
real    0m28.778s  
user    0m4.281s  
sys     0m24.507s
```



Q & A

